

# Jungfrau-Photoncounter

Florian Warg, Jonas Schenke, Sebastian Benner  
Dresden, January 31 2017


## Outline

- Task
- Implementation
  - Concept
  - Optimization Concepts
  - Results

## Task

- Conversion  
charge  $\rightarrow$  number of photons
- Module:
  - 2000 fps
  - 16 Bit / Pixel
  - $\sim 2$  GB/s

### Jungfrau-Sensor



## Task


- 3 gain stages

- Offset – pedestal maps

$$G_0: C_{\text{corr}}[x,y] = D_{\text{ADC}}[x,y] - P_0[x,y]$$

$$G_1: C_{\text{corr}}[x,y] = P_1[x,y] - D_{\text{ADC}}[x,y]$$

$$G_2: C_{\text{corr}}[x,y] = P_2[x,y] - D_{\text{ADC}}[x,y]$$



## Task

- Gain Maps

$$G_0: E_{\text{cal}}[x,y] = C_{\text{corr}}[x,y] * G_0[x,y]$$

$$G_1: E_{\text{cal}}[x,y] = C_{\text{corr}}[x,y] * G_1[x,y]$$



$$G_2: E_{\text{cal}}[x,y] = C_{\text{corr}}[x,y] * G_2[x,y]$$

- Data

- Offset removed

- Multiplied by gain map

- $NOP[x,y] = E_{\text{cal}}[x,y]/E_{\text{Beam}}$




## Implementation - Concept

### Upload function

- uploads data from buffer to GPU
- asynchronous kernel calculation
- creates CUDA callback

### Ringbuffer



### Download function

- copy data
- fill ringbuffer

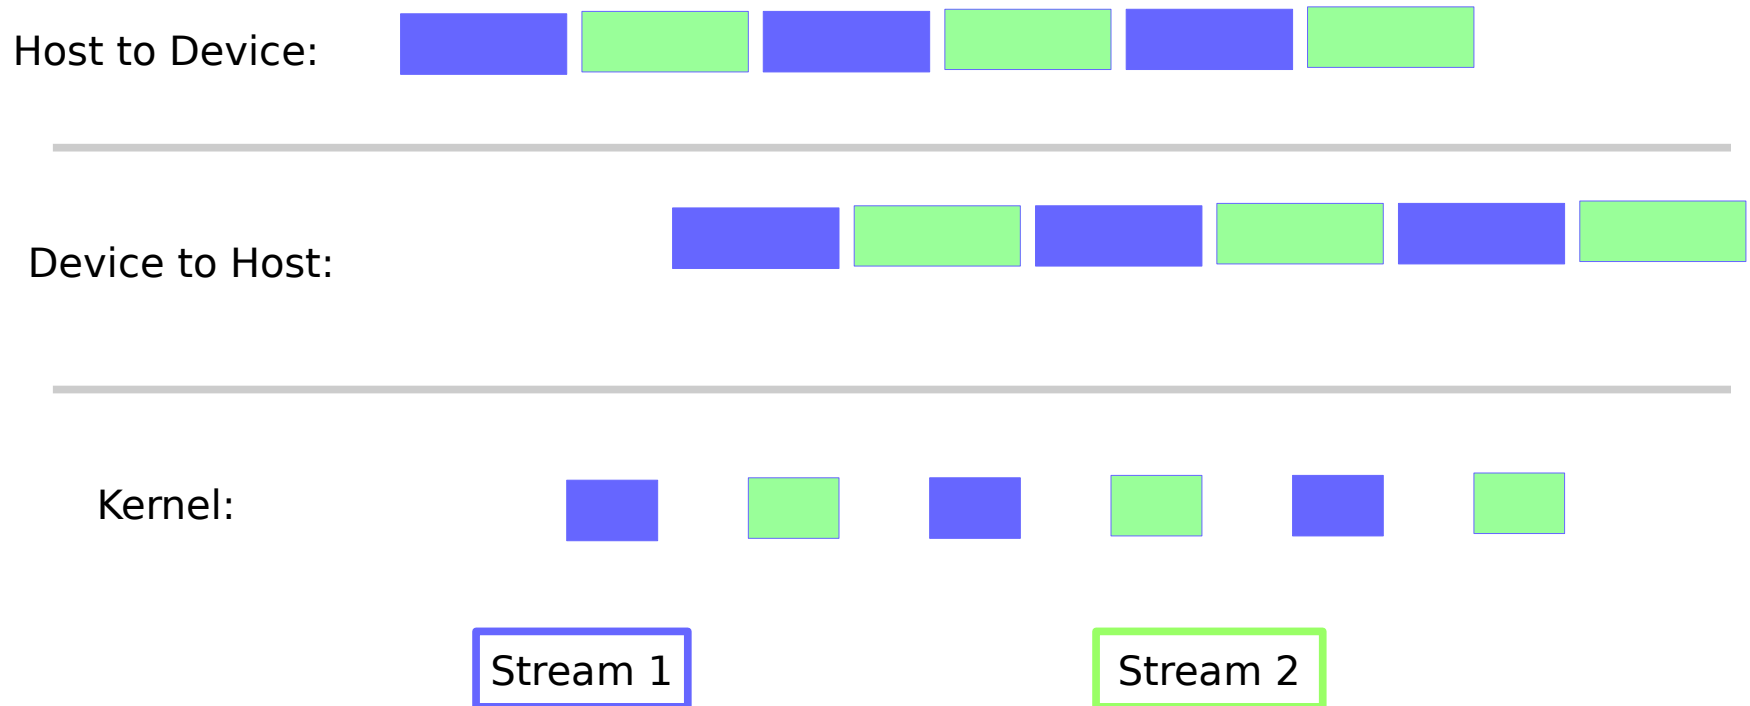
Vector  
<DeviceData>

---

DeviceData = streaminformationen; device- & host-pointer for data-,  
pedestal- and gainmaps; processing state

## Implementation - Concept

### Optimal Timeline:



## Implementation - Optimization Concepts

- Split frame blocks onto multiple devices
- Calculate multiple frames in one kernel
- Use streams:
  - Use pinned memory
  - Asynchronous data transfer
  - Asynchronous kernel call



## Implementation - Optimization Concepts

Further improvements (fine tuning)


- Find optimal number of frames
- Find optimal proportion streams / device
- Remove paged memory
- Sequential reads inside of the kernel
- Switch X and Y when calling kernel
  - $\text{kernel}\langle X, Y, \dots \rangle(\dots) \rightarrow \text{kernel}\langle Y, X, \dots \rangle(\dots)$

## Implementation - Optimization Concepts

- Current state:
  - Working (without pedestal calibration)
  - Multiple GPUs with multiple streams
  - Shared memory


## Implementation - Result

Current state:



# Implementation - Result

## Multiple GPUs:



# Implementation - Result

## Bandwidth Host to Device:

Name	Start Time	Duration	Size	Throughput
Memcpy HtoD [sync]	1,07 s	2,823 ms	12,583 MB	4,458 GB/s
Memcpy HtoD [sync]	1,073 s	110,202 µs	3,146 MB	7,669 GB/s
Memcpy HtoD [sync]	1,797 s	2,812 ms	12,583 MB	4,474 GB/s
Memcpy HtoD [sync]	1,8 s	112,219 µs	3,146 MB	7,631 GB/s
Memcpy HtoD [async]	3,724 s	84,735 ms	1,049 GB	12,375 GB/s
Memcpy HtoD [async]	4,659 s	84,727 ms	1,049 GB	12,376 GB/s
Memcpy HtoD [async]	8,206 s	84,717 ms	1,049 GB	12,377 GB/s
Memcpy HtoD [async]	9,216 s	84,737 ms	1,049 GB	12,374 GB/s
Memcpy HtoD [async]	12,574 s	84,735 ms	1,049 GB	12,375 GB/s
Memcpy HtoD [async]	13,585 s	84,74 ms	1,049 GB	12,374 GB/s
Memcpy HtoD [async]	16,792 s	84,735 ms	1,049 GB	12,375 GB/s
Memcpy HtoD [async]	17,801 s	84,71 ms	1,049 GB	12,378 GB/s
Memcpy HtoD [async]	21,162 s	84,707 ms	1,049 GB	12,379 GB/s
Memcpy HtoD [async]	22,174 s	84,708 ms	1,049 GB	12,379 GB/s

## Implementation - Result

### Bandwidth Device to Host:

Name	Start Time	Duration	Size	Throughput
Memcpy DtoH [async]	3,834 s	84,896 ms	1,049 GB	12,351 GB/s
Memcpy DtoH [async]	4,768 s	84,606 ms	1,049 GB	12,394 GB/s
Memcpy DtoH [async]	8,315 s	84,703 ms	1,049 GB	12,379 GB/s
Memcpy DtoH [async]	9,325 s	84,707 ms	1,049 GB	12,379 GB/s
Memcpy DtoH [async]	12,683 s	84,723 ms	1,049 GB	12,377 GB/s
Memcpy DtoH [async]	13,694 s	84,756 ms	1,049 GB	12,372 GB/s
Memcpy DtoH [async]	16,902 s	84,734 ms	1,049 GB	12,375 GB/s
Memcpy DtoH [async]	17,911 s	84,727 ms	1,049 GB	12,376 GB/s
Memcpy DtoH [async]	21,272 s	84,711 ms	1,049 GB	12,378 GB/s
Memcpy DtoH [async]	22,283 s	84,717 ms	1,049 GB	12,377 GB/s

# Implementation - Result

## Kernel:

Name	Start Time	Duration	Grid Size	Block Size	Regs
calculate(unsigned int, unsigned short*, double*, unsigned short*, unsigned int, unsigned short*)	3,809 s	24,899 ms	[1024,1,1]	[512,1,1]	32
calculate(unsigned int, unsigned short*, double*, unsigned short*, unsigned int, unsigned short*)	4,744 s	24,793 ms	[1024,1,1]	[512,1,1]	32
calculate(unsigned int, unsigned short*, double*, unsigned short*, unsigned int, unsigned short*)	8,29 s	24,811 ms	[1024,1,1]	[512,1,1]	32
calculate(unsigned int, unsigned short*, double*, unsigned short*, unsigned int, unsigned short*)	9,301 s	24,736 ms	[1024,1,1]	[512,1,1]	32
calculate(unsigned int, unsigned short*, double*, unsigned short*, unsigned int, unsigned short*)	12,659 s	24,81 ms	[1024,1,1]	[512,1,1]	32
calculate(unsigned int, unsigned short*, double*, unsigned short*, unsigned int, unsigned short*)	13,67 s	24,905 ms	[1024,1,1]	[512,1,1]	32
calculate(unsigned int, unsigned short*, double*, unsigned short*, unsigned int, unsigned short*)	16,877 s	24,791 ms	[1024,1,1]	[512,1,1]	32
calculate(unsigned int, unsigned short*, double*, unsigned short*, unsigned int, unsigned short*)	17,886 s	24,776 ms	[1024,1,1]	[512,1,1]	32
calculate(unsigned int, unsigned short*, double*, unsigned short*, unsigned int, unsigned short*)	21,247 s	24,821 ms	[1024,1,1]	[512,1,1]	32
calculate(unsigned int, unsigned short*, double*, unsigned short*, unsigned int, unsigned short*)	22,258 s	24,812 ms	[1024,1,1]	[512,1,1]	32

## Implementation - Result

Speedup Shared Memory:


- Global Memory:  $\sim 30$  ms
- Shared Memory:  $\sim 24.7$  ms  
→ Speedup:  $\sim 1.215$




## Implementation - Result

Output:

Result



Reference



# Questions?